Jincheng.Qian

SECURITY ENGINEERING · AI SAFETY

WHO I AM

Evidence before conclusions

A claim is only worth what a rerun can reproduce. I would rather say "unverified" than sound certain.

Skeptic by default — including of myself

I keep a second, independent line of thinking and let it attack the first one. If my own work survives that, I'll stand behind it.

Decisions are worth more than options

I don't collect choices, I commit to one and own the outcome. "We could still go back" reads to me as not having decided.

I know where I over-reach

I catch myself rationalising when a decision is uncomfortable. Naming it early is the only fix I've found that works.

Independent on purpose

I don't take on responsibilities beyond what I can carry well. Sharp work beats wide work, and I'd rather go deep on one thing than wide on five.

WHAT DRAWS ME IN

Logic & deduction

Tracing a chain of reasoning to the place where it quietly breaks.

Thinking across fields

The most interesting problems live in the gaps — where security meets behaviour, where systems meet people.

Universe physics

Not as a career, as a way to keep a sense of scale. It is where my sense of proportion comes from.

Craft & aesthetics

Cold palettes, deep space, engineered systems. I care about how something is built, not only what it does.

CURRENTLY

Building toward AI security

Network security and reverse engineering first; AI safety is the direction I'm heading in.

Working independently

Self-directed research, built and verified end to end.

Open to what fits

Research roles, security engineering, or anything that values rigour over volume.